

# HDS5230 Final Exam - programming - Miao Cai

May 12, 2019

## 1 HDS 5230 High Performance Computing

### 1.1 Final Exam - Programming Part

Miao Cai [miao.cai@slu.edu](mailto:miao.cai@slu.edu)

## 2 Introduction

The big goal is to use the provided dataset on health insurance charges to create a model that predicts charges as accurately as possible, based on the patient traits of age, sex, bmi, children, smoker, and region. As you generate this model, you should perform and document initial data quality checks, exploratory data analysis, and all of the models you try to fit.

## 3 Methods summary

All the data cleaning, visualization, and modeling were conducted in Python, and this reported was wrote in jupyternotebook. The Python session and package version information is shown below.

```
In [1]: import os
import sys
import pathlib
from tableone import TableOne
import pandas as pd
import numpy as np
import seaborn as sns
import h2o
from h2o.estimators.glm import H2OGeneralizedLinearEstimator

print(sys.version)
print("Pandas version: {}".format(pd.__version__))
print("Numpy version:{}".format(np.__version__))
print("Seaborn version:{}".format(sns.__version__))
print("h2o version:{}".format(h2o.__version__))
print("My working directory:\n" + os.getcwd())
```

```

3.7.1 (default, Dec 14 2018, 13:28:58)
[Clang 4.0.1 (tags/RELEASE_401/final)]
Pandas version: 0.23.4
Numpy version:1.15.4
Seaborn version:0.9.0
h2o version:3.22.1.3
My working directory:
/Users/miaocai/Dropbox/@2018 SPRING HDS5230 High performance computing/HDS5230Homework/Final e

```

### 3.0.1 Loss function

I pick the loss function as the mean absolute error (MAE). I chose this loss function since the outcome variable charges are highly right-skewed. Using the most commonly used mean square error (MSE) will not be as robust as MAE since the error will be squared, which is strongly influenced by outliers.

### 3.0.2 Initiate h2o and read data

```
In [2]: h2o.init(ip='localhost', port=54321, nthreads=-1, max_mem_size='2G')
```

```

Checking whether there is an H2O instance running at http://localhost:54321... not found.
Attempting to start a local H2O server...

```

```

Java Version: openjdk version "11.0.2" 2019-01-15; OpenJDK Runtime Environment 18.9 (build 1
Starting server from /Users/miaocai/anaconda3/lib/python3.7/site-packages/h2o/backend/bin/h2
Ice root: /var/folders/ng/t3l4gg0s0l398768705lv8vm0000gn/T/tmp8sc0eobc
JVM stdout: /var/folders/ng/t3l4gg0s0l398768705lv8vm0000gn/T/tmp8sc0eobc/h2o_miaocai_started
JVM stderr: /var/folders/ng/t3l4gg0s0l398768705lv8vm0000gn/T/tmp8sc0eobc/h2o_miaocai_started
Server is running at http://127.0.0.1:54321

```

```
Connecting to H2O server at http://127.0.0.1:54321... successful.
```

```
Warning: Your H2O cluster version is too old (3 months and 16 days)! Please download and install
```

```

-----
H2O cluster uptime:          01 secs
H2O cluster timezone:       America/Chicago
H2O data parsing timezone:   UTC
H2O cluster version:        3.22.1.3
H2O cluster version age:    3 months and 16 days !!!
H2O cluster name:           H2O_from_python_miaocai_k7904q
H2O cluster total nodes:    1
H2O cluster free memory:    2 Gb
H2O cluster total cores:    8
H2O cluster allowed cores:  8
H2O cluster status:         accepting new members, healthy
H2O connection url:         http://127.0.0.1:54321
H2O connection proxy:
H2O internal security:      False
H2O API Extensions:         XGBoost, Algos, AutoML, Core V3, Core V4

```

Python version: 3.7.1 final

---

```
In [3]: d0 = h2o.import_file("insurance.csv")
        d0
```

Parse progress: || 100%

Out[3]:

```
In [4]: print("The shape of the DataFrame is:")
        print(d0.shape)
        list(zip(d0.nacnt(), d0.names))
```

The shape of the DataFrame is:  
(1338, 7)

```
Out[4]: [(0.0, 'age'),
          (0.0, 'sex'),
          (0.0, 'bmi'),
          (0.0, 'children'),
          (0.0, 'smoker'),
          (0.0, 'region'),
          (0.0, 'charges')]
```

After reading the .csv file by using h2o, I found that there are no missing values in the data, which is great!

### 3.0.3 Split into train/validation/testing splits

We then need to convert some categorical variables into factors, and split the data into train, test, and validation sets.

```
In [5]: d0[['sex', 'smoker', 'region']] = d0[['sex', 'smoker', 'region']].asfactor()
        dtrain, dtest, dvalid = \
            d0.split_frame([0.7, 0.15], seed = 666)
```

## 4 Results

### 4.1 Summary statistics of model input variables

```
In [6]: d = pd.read_csv("insurance.csv")
        col_types = d.dtypes.to_dict()
        col_types['age'] = 'float64'
        d = pd.read_csv("insurance.csv", dtype=col_types)
```

```

overall_table = TableOne(
    d, columns = ['age', 'bmi', 'children', 'smoker', 'region'],
    categorical = ['children', 'smoker', 'region'],
    groupby = 'sex', label_suffix=True, pval = True)
overall_table

```

```

Out[6]:

```

		Grouped by sex				
variable	level	isnull	female	male	pval	
n			662	676		
age, mean (SD)		0	39.5 (14.1)	38.9 (14.1)	0.446	Two Sample T
bmi, mean (SD)		0	30.4 (6.0)	30.9 (6.1)	0.090	Two Sample T
children, n (%)	0	0	289 (43.7)	285 (42.2)	0.981	Chi-s
	1		158 (23.9)	166 (24.6)		
	2		119 (18.0)	121 (17.9)		
	3		77 (11.6)	80 (11.8)		
	4		11 (1.7)	14 (2.1)		
	5		8 (1.2)	10 (1.5)		
smoker, n (%)	no	0	547 (82.6)	517 (76.5)	0.007	Chi-s
	yes		115 (17.4)	159 (23.5)		
region, n (%)	northeast	0	161 (24.3)	163 (24.1)	0.933	Chi-s
	northwest		164 (24.8)	161 (23.8)		
	southeast		175 (26.4)	189 (28.0)		
	southwest		162 (24.5)	163 (24.1)		

[1] Warning, Hartigan's Dip Test reports possible multimodal distributions for: age.  
[2] Warning, test for normality reports non-normal distributions for: age.

## 4.2 Descriptive analysis of the outcome variable charges

```

In [7]: sns.distplot(d['charges'], hist=True, kde=True,
    bins=int(180/5), color = 'skyblue',
    hist_kws={'edgecolor':'grey'},
    kde_kws={'linewidth': 2})

```

```


/Users/miaocai/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning:
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

```

```


Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1badc160>

```



```
In [8]: sns.boxplot(d['charges'], color = 'skyblue')
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1c00b978>
```



```
In [9]: d['charges'].describe()
```

```
Out[9]: count      1338.000000
        mean       13270.422265
        std        12110.011237
        min        1121.873900
        25%        4740.287150
        50%        9382.033000
        75%        16639.912515
        max        63770.428010
        Name: charges, dtype: float64
```

### 4.3 Cross-validation

Here I use the `H2OGeneralizedLinearEstimator` module in `h2o` package to perform cross-validation. I set the number of cross-validation as 5 by `nfolds=5`, in which 6 models will be built. The first 5 models (cross-validation models) are built on 80% of the training data, and a different 20% is held out for each of the 5 models. Then the main model is built on 100% of the training data.

```
In [10]: target_var = 'charges'
        input_var = ['age', 'sex', 'bmi',
                     'children', 'smoker', 'region']
        glm_mod0 = \
            H2OGeneralizedLinearEstimator(model_id = 'd_glm_0',
                                           family = 'gaussian',
                                           nfolds = 5,
                                           seed = 123)

        glm_mod0.train(x = input_var,
                       y = target_var,
                       training_frame = d0)

        glm_mod0.show()
```

```
glm Model Build progress: || 100%
```

```
Model Details
```

```
=====
```

```
H2OGeneralizedLinearEstimator : Generalized Linear Modeling
```

```
Model Key: d_glm_0
```

```
ModelMetricsRegressionGLM: glm
```

```
** Reported on train data. **
```

```
MSE: 125728309.44330035
```

```
RMSE: 11212.863570172445
```

```
MAE: 8471.397911142334
```

RMSLE: 0.9331157151413909  
 R<sup>2</sup>: 0.14203674139265943  
 Mean Residual Deviance: 125728309.44330035  
 Null degrees of freedom: 1337  
 Residual degrees of freedom: 1326  
 Null deviance: 196074221532.74988  
 Residual deviance: 168224478035.13586  
 AIC: 28776.289622947068

ModelMetricsRegressionGLM: glm  
 \*\* Reported on cross-validation data. \*\*

MSE: 129830749.84711617  
 RMSE: 11394.32972346843  
 MAE: 8600.97548619905  
 RMSLE: 0.9464119017196216  
 R<sup>2</sup>: 0.11404190751089327  
 Mean Residual Deviance: 129830749.84711617  
 Null degrees of freedom: 1337  
 Residual degrees of freedom: 1326  
 Null deviance: 196632689167.88873  
 Residual deviance: 173713543295.44144  
 AIC: 28819.25062652632  
 Cross-Validation Metrics Summary:

	mean	sd	cv_1_valid	cv_2_valid	cv_3_valid	cv_4_valid
-----	-----	-----	-----	-----	-----	-----
mae	8592.55	295.311	8582.31	7818.53	8785.11	8785.11
mean_residual_deviance	1.29276e+08	1.25685e+07	1.22953e+08	1.02176e+08	1.53851e+08	1.53851e+08
mse	1.29276e+08	1.25685e+07	1.22953e+08	1.02176e+08	1.53851e+08	1.53851e+08
null_deviance	3.93265e+10	4.53807e+09	3.72839e+10	2.96943e+10	4.90307e+10	3.72839e+10
r2	0.110559	0.00589982	0.117462	0.114765	0.100491	0.100491
residual_deviance	3.47427e+10	4.17128e+09	3.27055e+10	2.6157e+10	4.38475e+10	3.27055e+10
rmse	11342.6	557.598	11088.4	10108.2	12403.7	11088.4
rmsle	0.945845	0.0343231	0.981383	0.914074	0.884154	0.884154

Scoring History:

	timestamp	duration	iterations	negative_log_likelihood	objective
--	-----	-----	-----	-----	-----
	2019-05-12 04:09:36	0.000 sec	0	1.96074e+11	1.46543e+08

## 4.4 Machine learning

### 4.4.1 Regularized regression

```
In [11]: from h2o.grid.grid_search import H2OGridSearch
```

```
alpha_values = {'alpha': [0, 0.25, 0.5, 0.75, 1]}
glm_mod2 = H2OGridSearch(
    H2OGeneralizedLinearEstimator(family = 'gaussian',
                                   lambda_search = True,
                                   seed = 123), \
    hyper_params = alpha_values)

glm_mod2.train(x = input_var,
               y = target_var,
               training_frame = dtrain,
               validation_frame = dvalid)

glm_mod2.show()
```

```
glm Grid Build progress: || 100%
```

	alpha	model_ids \
0	[1.0]	Grid_GLM_py_4_sid_b6da_model_python_1557652172496_1_model_5
1	[0.75]	Grid_GLM_py_4_sid_b6da_model_python_1557652172496_1_model_4
2	[0.5]	Grid_GLM_py_4_sid_b6da_model_python_1557652172496_1_model_3
3	[0.25]	Grid_GLM_py_4_sid_b6da_model_python_1557652172496_1_model_2
4	[0.0]	Grid_GLM_py_4_sid_b6da_model_python_1557652172496_1_model_1

	residual_deviance
0	6.232294045221491E9
1	2.7448416102541004E10
2	2.7451841045226257E10
3	2.7452983289290577E10
4	2.7453233838145943E10

```
In [12]: glm_best = glm_mod2.get_grid()[0]
         glm_best._model_json['output']['coefficients_table'].as_data_frame()
```

```
Out[12]:
```

	names	coefficients	standardized_coefficients
0	Intercept	-7206.776563	13492.999531
1	region.northeast	83.984427	83.984427
2	region.northwest	111.459665	111.459665
3	region.southeast	-192.376821	-192.376821
4	region.southwest	-70.256674	-70.256674
5	smoker.no	-5020.292788	-5020.292788
6	smoker.yes	18846.423625	18846.423625
7	sex.female	147.142857	147.142857
8	sex.male	-142.895434	-142.895434



9	age	255.041295	3604.495432
10	bmi	329.352274	2038.979286
11	children	477.291820	563.040100

```
In [13]: glm_best._model_json['output']['training_metrics']
```

```
ModelMetricsRegressionGLM: glm
** Reported on train data. **
```

```
MSE: 36244089.05252515
RMSE: 6020.306391914381
MAE: 4175.590832941357
RMSLE: NaN
R^2: 0.7557379707834732
Mean Residual Deviance: 36244089.05252515
Null degrees of freedom: 955
Residual degrees of freedom: 944
Null deviance: 141853194478.7002
Residual deviance: 34649349134.21404
AIC: 19378.942718761107
```

```
Out[13]:
```

```
In [14]: glm_best._model_json['output']['validation_metrics']
```

```
ModelMetricsRegressionGLM: glm
** Reported on validation data. **
```

```
MSE: 30107700.701553095
RMSE: 5487.04845081152
MAE: 3827.8094066015087
RMSLE: 0.5274009774363978
R^2: 0.7699240420817869
Mean Residual Deviance: 30107700.701553095
Null degrees of freedom: 206
Residual degrees of freedom: 195
Null deviance: 27453554523.254715
Residual deviance: 6232294045.221491
AIC: 4178.040900338396
```

```
Out[14]:
```

#### 4.4.2 Auto-ML

```
In [15]: from h2o.automl.autoh2o import H2OAutoML
```

```

automl_1 = H2OAutoML(max_runtime_secs = 60)
automl_1.train(x = input_var,
               y = target_var,
               training_frame = dtrain,
               validation_frame = dvalid)
automl_1.leaderboard

```

AutoML progress: || 100%

Out[15]:

In [16]: automl\_1.leader

Model Details

=====

H2OXGBoostEstimator : XGBoost  
Model Key: XGBoost\_3\_AutoML\_20190512\_040937

ModelMetricsRegression: xgboost

\*\* Reported on train data. \*\*

MSE: 12188391.73287312  
RMSE: 3491.187725240956  
MAE: 1819.1937206060816  
RMSLE: 0.3021820933709707  
Mean Residual Deviance: 12188391.73287312

ModelMetricsRegression: xgboost

\*\* Reported on validation data. \*\*

MSE: 19380121.92359663  
RMSE: 4402.285988392466  
MAE: 2431.2812163864355  
RMSLE: 0.40985324039544335  
Mean Residual Deviance: 19380121.92359663

ModelMetricsRegression: xgboost

\*\* Reported on cross-validation data. \*\*

MSE: 19998514.578738146  
RMSE: 4471.969876769984  
MAE: 2408.4270772894056  
RMSLE: 0.41859298826412966  
Mean Residual Deviance: 19998514.578738146  
Cross-Validation Metrics Summary:

	mean	sd	cv_1_valid	cv_2_valid	cv_3_valid	cv_4_valid
-----	-----	-----	-----	-----	-----	-----
mae	2408.46	82.2232	2379.37	2247.67	2338.2	2500.0
mean_residual_deviance	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
mse	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
r2	0.862954	0.0128701	0.876166	0.829337	0.880435	0.880435
residual_deviance	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
rmse	4465.8	163.311	4657.31	4508.39	4190.79	4200.0
rmsle	0.417898	0.0173632	0.393145	0.448762	0.416373	0.416373

#### Scoring History:

timestamp	duration	number_of_trees	training_rmse	training_mae	training_r2
--	-----	-----	-----	-----	-----
2019-05-12 04:09:49	1.424 sec	0	18177.7	13492.5	3.304
2019-05-12 04:09:49	1.437 sec	5	14489.6	10517.1	2.099
2019-05-12 04:09:49	1.445 sec	10	11629.2	8212.83	1.352
2019-05-12 04:09:49	1.456 sec	15	9524.79	6461.63	9.072
2019-05-12 04:09:49	1.465 sec	20	7934.06	5154.68	6.294
2019-05-12 04:09:49	1.475 sec	25	6704.42	4165.21	4.494
2019-05-12 04:09:49	1.487 sec	30	5943.49	3498.78	3.532
2019-05-12 04:09:49	1.499 sec	35	5275.12	2966.53	2.782
2019-05-12 04:09:49	1.512 sec	40	4770.26	2572.71	2.275
2019-05-12 04:09:49	1.526 sec	45	4406.44	2309.07	1.941
2019-05-12 04:09:49	1.542 sec	50	4171.37	2145.19	1.740
2019-05-12 04:09:49	1.559 sec	55	3978.81	2038	1.583
2019-05-12 04:09:49	1.577 sec	60	3815.2	1950.83	1.455
2019-05-12 04:09:49	1.596 sec	65	3698.73	1890.03	1.368
2019-05-12 04:09:50	1.617 sec	70	3620.38	1864.41	1.310
2019-05-12 04:09:50	1.639 sec	75	3548.99	1837.89	1.259
2019-05-12 04:09:50	1.663 sec	80	3501.17	1824.03	1.225
2019-05-12 04:09:50	1.683 sec	81	3491.19	1819.19	1.218

#### Variable Importances:

variable	relative_importance	scaled_importance	percentage
-----	-----	-----	-----
smoker.no	4.10118e+11	1	0.499855
smoker.yes	1.4899e+11	0.363285	0.18159
bmi	1.35328e+11	0.329973	0.164939
age	1.02284e+11	0.2494	0.124664
children	1.0496e+10	0.0255926	0.0127926
region.northeast	4.24613e+09	0.0103534	0.00517522
sex.female	2.46772e+09	0.00601708	0.00300767
region.southeast	1.95259e+09	0.00476104	0.00237983

sex.male	1.7502e+09	0.00426755	0.00213316
region.southwest	1.43316e+09	0.00349449	0.00174674
region.northwest	1.40847e+09	0.00343429	0.00171665

Out[16]:

```
In [17]: autom1_1.leader.cross_validation_metrics_summary()
```

Cross-Validation Metrics Summary:

	mean	sd	cv_1_valid	cv_2_valid	cv_3_valid	cv_4_valid
mae	2408.46	82.2232	2379.37	2247.67	2338.2	2500.0
mean_residual_deviance	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
mse	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
r2	0.862954	0.0128701	0.876166	0.829337	0.880435	0.880435
residual_deviance	1.99967e+07	1.45611e+06	2.16905e+07	2.03256e+07	1.75627e+07	1.75627e+07
rmse	4465.8	163.311	4657.31	4508.39	4190.79	4200.0
rmsle	0.417898	0.0173632	0.393145	0.448762	0.416373	0.416373

Out[17]:

## 4.5 Neural networks using keras and tensorflow

```
In [18]: from keras import models
         from keras import layers
         from keras import optimizers
         from sklearn import preprocessing
         from sklearn.model_selection import train_test_split
         import matplotlib.pyplot as plt

d = pd.read_csv("insurance.csv")
col_types = d.dtypes.to_dict()
col_types['age'] = 'float64'
d = pd.read_csv("insurance.csv", dtype=col_types)

d_dummy = pd.get_dummies(d, drop_first = True)
y = d_dummy['charges'].values
X = d_dummy.drop('charges', axis = 1).values

X_train0, X_test0 = train_test_split(
    X, train_size = 0.8, random_state = 123)
y_train0, y_test0 = train_test_split(
    y, train_size = 0.8, random_state = 123)
X_train1, X_valid1 = train_test_split(
    X, train_size = 0.75, random_state = 123)
```

```

y_train1, y_valid1 = train_test_split(
    y, train_size = 0.75, random_state = 123)

# logistic regression
def nn_model1():
    model = models.Sequential()
    model.add(layers.Dense(64, activation = 'relu',
                           input_dim = X_train1.shape[1]))
    model.add(layers.Dense(1,
                           activation = 'linear'))
    model.compile(optimizer = 'rmsprop',
                  loss = 'mse',
                  metrics = ['mae'])
    return model

reg0 = nn_model1()
reg0.fit(X_train1, y_train1, epochs = 200,
        validation_data = [X_valid1, y_valid1],
        verbose = 0)

```

Using TensorFlow backend.

/Users/miaocai/anaconda3/lib/python3.7/site-packages/sklearn/model\_selection/\_split.py:2179: FutureWarning)

FutureWarning)

WARNING:tensorflow:From /Users/miaocai/anaconda3/lib/python3.7/site-packages/tensorflow/python/

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /Users/miaocai/anaconda3/lib/python3.7/site-packages/tensorflow/python/

Instructions for updating:

Use tf.cast instead.

Out[18]: <keras.callbacks.History at 0x1a2d078f28>

In [19]: reg0.history.history['loss'][199]

Out[19]: 125999988.09172483

In [20]: reg0.history.history['val\_loss'][199]


Out[20]: 133231693.99402985

In [21]: reg0.history.history['val\_mean\_absolute\_error'][199]

Out[21]: 9078.39338123834


In [22]: plt.plot(reg0.history.history['loss'], color = 'red')  
plt.plot(reg0.history.history['val\_loss'], color = 'blue')

Out[22]: [<matplotlib.lines.Line2D at 0x1a2d822ba8>]



```
In [23]: plt.plot(reg0.history.history['mean_absolute_error'], color = 'red')
         plt.plot(reg0.history.history['val_mean_absolute_error'], color = 'blue')
```

```
Out[23]: [<matplotlib.lines.Line2D at 0x1a2d7dbda0>]
```



```

In [24]: def nn_model2():
        model = models.Sequential()
        model.add(layers.Dense(64, activation = 'relu',
                                input_dim = X_train1.shape[1]))
        model.add(layers.Dropout(0.5))
        model.add(layers.Dense(64, activation = 'relu'))
        model.add(layers.Dropout(0.5))
        model.add(layers.Dense(1, activation = 'linear'))
        model.compile(optimizer = 'rmsprop',
                      loss = 'mse',
                      metrics = ['mae'])
        return model

reg2 = nn_model2()
reg2.fit(X_train1, y_train1, epochs = 200,
        validation_data = [X_valid1, y_valid1],
        verbose = 0)

```

WARNING:tensorflow:From /Users/miaocai/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:1445: tf.nn.conv2d is deprecated and will be removed in a future version. Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.


Out [24]: <keras.callbacks.History at 0x1a2d8c7048>

```

In [25]: plt.plot(reg2.history.history['loss'], color = 'red')
        plt.plot(reg2.history.history['val_loss'], color = 'blue')


```

Out [25]: [<matplotlib.lines.Line2D at 0x1a2dde9eb8>]



```
In [26]: plt.plot(reg2.history.history['mean_absolute_error'], color = 'red')
         plt.plot(reg2.history.history['val_mean_absolute_error'], color = 'blue')
```

```
Out[26]: [<matplotlib.lines.Line2D at 0x1a2deb2550>]
```




```
In [27]: from keras.regularizers import l1
```

```
def nn_model3(l1_penalty):
    model = models.Sequential()
    model.add(layers.Dense(64, activation = 'relu',
                           kernel_regularizer =
                               l1(l1_penalty),
                           input_dim = X_train1.shape[1]))
    model.add(layers.Dropout(0.5))
    model.add(layers.Dense(64, activation = 'relu',
                           kernel_regularizer = l1(l1_penalty)))
    model.add(layers.Dropout(0.5))
    model.add(layers.Dense(1, activation = 'linear'))
    model.compile(optimizer = 'rmsprop',
                  loss = 'mse',
                  metrics = ['mae'])
    return model
```




```
In [28]: reg3 = nn_model3(0.1)
reg3.fit(X_train1, y_train1, epochs = 200,
        validation_data = [X_valid1, y_valid1],
        verbose = 0)
plt.plot(reg3.history.history['loss'], color = 'red')
plt.plot(reg3.history.history['val_loss'], color = 'blue')
```

Out[28]: [




```
In [29]: plt.plot(reg3.history.history['mean_absolute_error'], color = 'red')
plt.plot(reg3.history.history['val_mean_absolute_error'], color = 'blue')
```

Out[29]: [




```
In [30]: reg3 = nn_model3(0.2)
         reg3.fit(X_train1, y_train1, epochs = 200,
                 validation_data = [X_valid1, y_valid1],
                 verbose = 0)
         plt.plot(reg3.history.history['loss'], color = 'red')
         plt.plot(reg3.history.history['val_loss'], color = 'blue')
```

```
Out[30]: [<matplotlib.lines.Line2D at 0x1a2eab65f8>]
```




```
In [31]: plt.plot(reg3.history.history['mean_absolute_error'], color = 'red')
         plt.plot(reg3.history.history['val_mean_absolute_error'], color = 'blue')
```

```
Out[31]: [<matplotlib.lines.Line2D at 0x1a2eb15cf8>]
```




```
In [32]: reg3 = nn_model3(0.3)
reg3.fit(X_train1, y_train1, epochs = 200,
        validation_data = [X_valid1, y_valid1],
        verbose = 0)
plt.plot(reg3.history.history['loss'], color = 'red')
plt.plot(reg3.history.history['val_loss'], color = 'blue')
```

Out[32]: [



```
In [33]: plt.plot(reg3.history.history['mean_absolute_error'], color = 'red')
plt.plot(reg3.history.history['val_mean_absolute_error'], color = 'blue')
```

Out[33]: [



## 5 Conclusion

In [ ]:

In [ ]:

In [ ]: