

# HDS 5230 High Performance Computing - HW5

February 20, 2019

Author: Miao Cai

```
In [ ]: import pandas as pd
import os
print("My working directory:\n" + os.getcwd())
# os.chdir(r"C:\Users\evancarey\Dropbox\Work\SLU\Courses")
```

a. import only the first 100 rows of the patient.csv dataset using Python/pandas

```
In [22]: pt = pd.read_csv('healthcare2/Patient.csv', nrows = 100)
pt.loc[0:10,:]
```

```
Out[22]:
```

	PatientID	FirstName	LastName	State	ZipCode	DateOfBirth	Gender	\
0	1	Diana	Huddleston	WI	53186	1962-02-27	female	
1	2	Marion	Poston	IL	60527	1859-09-11	male	
2	3	Sandra	Hamby	IL	60126	1946-02-15	female	
3	4	Mildred	Krehbiel	ID	83702	1979-07-27	female	
4	5	Abigail	Flores	PA	19131	1983-02-19	female	
5	6	Rusty	Thomas	AL	36107	NaN	male	
6	7	Robert	Alexander	CA	94539	1958-01-11	male	
7	8	Krista	Ward	WI	53219	1952-10-31	female	
8	9	Marti	Calabrese	MS	38801	1951-10-06	female	
9	10	Jeremy	Liu	CA	95526	1954-10-16	male	
10	11	Catherine	Tatum	MI	48213	1983-07-12	female	

  

	Race	Income
0	NaN	1076.167979
1	white	475.781094
2	white	30.747987
3	white	160.596425
4	?	NaN
5	black	171.378008
6	Missing	66.226314
7	black	15.078950
8	Missing	114.598911
9	white	1081.877157
10	Missing	35.058641

b. Examine the column names and the dtypes of the dataframe

```
In [23]: pt.columns
```

```
Out[23]: Index(['PatientID', 'FirstName', 'LastName', 'State', 'ZipCode', 'DateOfBirth',  
              'Gender', 'Race', 'Income'],  
              dtype='object')
```

```
In [24]: pt.dtypes
```

```
Out[24]: PatientID      int64  
         FirstName     object  
         LastName      object  
         State         object  
         ZipCode       int64  
         DateOfBirth   object  
         Gender        object  
         Race          object  
         Income        float64  
         dtype: object
```

c. Create a dict of columns names and types using the `to_dict()` method

```
In [25]: col_types = pt.dtypes.to_dict()  
         col_types
```

```
Out[25]: {'PatientID': dtype('int64'),  
          'FirstName': dtype('O'),  
          'LastName': dtype('O'),  
          'State': dtype('O'),  
          'ZipCode': dtype('int64'),  
          'DateOfBirth': dtype('O'),  
          'Gender': dtype('O'),  
          'Race': dtype('O'),  
          'Income': dtype('float64')}
```

d. Decide which columns you can compress by specifying a smaller dtype. For example, the default dtype of an integer is `int64`, but you may be able to fit that integer data into the dtype `int32`, or `int16`, or `uint16` (unsigned integer). It depends on the data! Consider turning the text data into categorical data. Try to make the dataframe as small as reasonably possible.

```
In [26]: pt.info(memory_usage='deep')  
         col_types['PatientID']='uint16'  
         col_types['FirstName']='category'  
         col_types['LastName']='category'  
         col_types['State']='category'  
         col_types['ZipCode']='uint16'  
         col_types['DateOfBirth']='category'  
         col_types['Gender']='category'  
         col_types['Race']='category'  
         col_types['Income']='float32'
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 9 columns):
PatientID      100 non-null int64
FirstName      100 non-null object
LastName       100 non-null object
State          100 non-null object
ZipCode        100 non-null int64
DateOfBirth    92 non-null object
Gender         98 non-null object
Race           97 non-null object
Income         92 non-null float64
dtypes: float64(1), int64(2), object(6)
memory usage: 38.8 KB

```

- e. Use the `memory_usage(deep=True)` dataframe method to calculate large your reduced file is.

```

In [27]: pt_reduced = pd.read_csv('healthcare2/Patient.csv', nrows = 100,
dtype=col_types)
pt_reduced.info(memory_usage='deep')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 9 columns):
PatientID      100 non-null uint16
FirstName      100 non-null category
LastName       100 non-null category
State          100 non-null category
ZipCode        100 non-null uint16
DateOfBirth    92 non-null category
Gender         98 non-null category
Race           97 non-null category
Income         92 non-null float32
dtypes: category(6), float32(1), uint16(2)
memory usage: 30.7 KB

```

- f. Import the `patient.csv` dataframe with default datatypes and calculate the `memory_usage(deep=True)`. How much smaller is your reduced dataframe than the full dataframe?

```

In [28]: pt1 = pd.read_csv('healthcare2/Patient.csv')
pt1.info(memory_usage='deep')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 9 columns):

```

```

PatientID      20000 non-null int64
FirstName      20000 non-null object
LastName       20000 non-null object
State          20000 non-null object
ZipCode        20000 non-null int64
DateOfBirth    19000 non-null object
Gender         19431 non-null object
Race           19144 non-null object
Income         18600 non-null float64
dtypes: float64(1), int64(2), object(6)
memory usage: 7.6 MB

```

```

In [29]: pt1_reduced = pd.read_csv('healthcare2/Patient.csv', dtype=col_types)
         pt1_reduced.info(memory_usage='deep')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 9 columns):
PatientID      20000 non-null uint16
FirstName      20000 non-null category
LastName       20000 non-null category
State          20000 non-null category
ZipCode        20000 non-null uint16
DateOfBirth    19000 non-null category
Gender         19431 non-null category
Race           19144 non-null category
Income         18600 non-null float32
dtypes: category(6), float32(1), uint16(2)
memory usage: 2.8 MB

```

g. Repeat a-f for the OutpatientVisit.csv file.

```

In [30]: op0 = pd.read_csv('healthcare2/OutpatientVisit.csv', nrows = 100)
         op0.loc[0:10,: ]

```

```

Out[30]:
   VisitID  StaffID  PatientID  VisitDate  ICD10_1  ICD10_2  ICD10_3  \
0         1        46         1  2013-08-10  E10621    K269      NaN
1         2        50         1  2013-12-02    K269  E10621      NaN
2         3        13         1  2014-06-29  E10621    K269      NaN
3         4        23         1  2014-09-19    K269  E10621      NaN
4         5         9         1  2015-05-29    K269  E10621      NaN
5         6        46         1  2016-05-07  E10621    K269      NaN
6         7         7         1  2016-10-07  E10621    K269      NaN
7         8        18         1  2016-11-07    K269  E10621      NaN
8         9        23         1  2017-01-14    K269  E10621      NaN
9        10         5         1  2017-01-29  E10621    K269      NaN
10        11         2         1  2017-06-29    K269  E10621      NaN

```

	ClinicCode
0	15
1	55
2	1
3	3
4	5
5	15
6	41
7	31
8	3
9	14
10	55

```
In [31]: op0.columns
```

```
Out[31]: Index(['VisitID', 'StaffID', 'PatientID', 'VisitDate', 'ICD10_1', 'ICD10_2',
               'ICD10_3', 'ClinicCode'],
              dtype='object')
```

```
In [32]: op0.dtypes
```

```
Out[32]: VisitID      int64
StaffID      int64
PatientID    int64
VisitDate    object
ICD10_1      object
ICD10_2      object
ICD10_3      float64
ClinicCode   int64
dtype: object
```

```
In [33]: col_types = op0.dtypes.to_dict()
col_types
```

```
Out[33]: {'VisitID': dtype('int64'),
          'StaffID': dtype('int64'),
          'PatientID': dtype('int64'),
          'VisitDate': dtype('O'),
          'ICD10_1': dtype('O'),
          'ICD10_2': dtype('O'),
          'ICD10_3': dtype('float64'),
          'ClinicCode': dtype('int64')}
```

```
In [34]: pt.info(memory_usage='deep')
col_types['VisitID']='uint16'
col_types['StaffID']='uint16'
col_types['PatientID']='uint16'
col_types['VisitDate']='category'
```

```
col_types['ICD10_1']='category'
col_types['ICD10_2']='category'
col_types['ICD10_3']='category'
col_types['ClinicCode']='uint16'
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 9 columns):
PatientID      100 non-null int64
FirstName      100 non-null object
LastName       100 non-null object
State          100 non-null object
ZipCode        100 non-null int64
DateOfBirth    92 non-null object
Gender         98 non-null object
Race           97 non-null object
Income         92 non-null float64
dtypes: float64(1), int64(2), object(6)
memory usage: 38.8 KB
```

```
In [35]: op0_reduced = pd.read_csv('healthcare2/OutpatientVisit.csv', nrows = 100,
                                   dtype=col_types)
        op0_reduced.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 8 columns):
VisitID        100 non-null uint16
StaffID        100 non-null uint16
PatientID      100 non-null uint16
VisitDate      100 non-null category
ICD10_1        100 non-null category
ICD10_2        63 non-null category
ICD10_3        0 non-null category
ClinicCode     100 non-null uint16
dtypes: category(4), uint16(4)
memory usage: 12.7 KB
```

```
In [36]: op = pd.read_csv('healthcare2/OutpatientVisit.csv')
        op.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 174690 entries, 0 to 174689
Data columns (total 8 columns):
VisitID        174690 non-null int64
StaffID        174690 non-null int64
PatientID      174690 non-null int64
```

```
VisitDate      173252 non-null object
ICD10_1         174690 non-null object
ICD10_2         59785 non-null object
ICD10_3         19362 non-null object
ClinicCode      174690 non-null int64
dtypes: int64(4), object(4)
memory usage: 39.7 MB
```

```
In [37]: op_reduced = pd.read_csv('healthcare2/OutpatientVisit.csv', dtype=col_types)
        op_reduced.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 174690 entries, 0 to 174689
Data columns (total 8 columns):
VisitID        174690 non-null uint16
StaffID        174690 non-null uint16
PatientID      174690 non-null uint16
VisitDate      173252 non-null category
ICD10_1        174690 non-null category
ICD10_2        59785 non-null category
ICD10_3        19362 non-null category
ClinicCode     174690 non-null uint16
dtypes: category(4), uint16(4)
memory usage: 3.7 MB
```